

# Profile directed meta-programming

William J. Bowman <wilbowma@ccs.neu.edu>,  
Swaha Miller <swamille@cisco.com>,  
and R. Kent Dybvig <dyb@cisco.com>

## Abstract

Many languages, such as ML, C++, Haskell, Java, and Scheme, provide powerful meta-programming facilities that help programmers create generic libraries, new language constructs, and domain-specific languages (DSLs). Meta-program manipulate source programs in ways reminiscent of a compiler. In some ways, meta-programs serve the same purpose as a compiler—to produce low-level code that is difficult to reason about from the high-level code that is easy to understand. With meta-programming increasingly prevalent, meta-programmers need some of the same tools and techniques compiler writers have been developing for decades.

Profile directed optimization is a compiler technique that uses sample data gathered at run-time to recompile and further optimize a program. The profile data can be more accurate than static heuristics normally used in a compiler and thus can lead to more optimized code. Today's compilers like .NET, LLVM, and GCC all provide some kind of profile directed optimizations.

This paper presents a system for using profile information to optimize programs via meta-programming. The system is implemented and used in a high-performance implementation of Scheme.

## 1. Introduction

Profile directed optimization is a standard compiler technique that uses sample data gathered at run-time to recompile and further optimize a program. If the sample data is representative of how the program is used in practice then this data can be more accurate than static heuristics and can lead to a more optimized program.

Compilers such as .NET, GCC, and LLVM do low level profile directed optimization such as reordering blocks, unrolling loops, and optimizing switch statements. These optimizations are often fragile because of their dependence on low level, compiler internal structures. Small changes to source code can imply large changes to low level representations. By bringing profile directed optimizations up to the source level, programmers can use specific knowledge of the problem domain, program, and the high-level code to optimize programs in ways a compiler can not.

A simple example is a conditional branching construct like Scheme's `cond`. `cond` takes an arbitrary number of clauses of the form `(lhs rhs)`, executing the first right-hand side whose left-hand side is true, or executing a final `else` clause if no left-hand

side is true. This construct essentially expands into a sequence of if/else expressions. However, if the programmer knows that each clause is mutually exclusive, it is beneficial to sort the clauses from most to least likely to succeed. In general, the compiler cannot prove such a property and must emit the clauses in the original order, so even traditional profile directed optimization cannot optimize `cond`.

This paper presents a system for doing profile directed meta-programming, including a workflow for using it with traditional low level profile directed optimizations. Section 2 presents the design of our system at a high level and how other macro systems could use it, section 3 presents several examples of Scheme macros that use this system, and section 4 discusses how this profiling system is implemented.

While Scheme is used in this paper, the same techniques should work in any language with sufficient meta-programming capabilities, such as Template Haskell, C++ Templates, or MacroML.

## 2. Design

This section presents the design of our profile system. We discuss the system at a high-level and sketch implementations for other macro systems. We discuss implementation details in section 4

### 2.1 Source and syntax objects

In Scheme, macros operate on *syntax objects*, and can run arbitrary Scheme code. To access profile information, all we require is a function that allows retrieving profile information from a syntax object. We added the function `profile-query-weight` to our Scheme implementation. Given a syntax object, `profile-query-weight` returns a number between 0 and 1, or false if there is no profile information associated with that piece of syntax.

### 2.2 Profile weight

We represent profile information as a floating point number between 0 and 1. Profile information is not stored as exact counts, but as execution frequency with respect to the most executed expression (referred to as 'percent of max'). If an expression `e1` is executed 1 time, and the most frequently executed expression `e10` is executed 10 times, then `(profile-query-weight e1)` returns .1, while `(profile-query-weight e10)` returns 1.

We use percent of max count in part because an exact execution count can be meaningless in some contexts. Consider an expression that is executed 5 times. We cannot know if this statement is executed frequently or not without some comparison.

We choose percent of max because this seems to give the best relative comparison. We considered comparing to the total number of expressions executed and the average number of times an expression is executed. In both cases, the results are distorted when there are a large number of expressions that are executed infrequently. In that case, a main loop might look infrequently executed if there are

Need to look at more related work before I can do this section justice

I need to more carefully read related work

many start up or shut down steps. By comparing to the most expensive expression, we have a relatively stable comparison of how expensive some expression is.

This relative information is not perfect. Loop unrolling can benefit from exact counts more than a weight. If we know a loop is executed exactly 5 times, unrolling it 5 times might make sense. If we know a loop is executed 20% of the max, we do not know if the loop is executed 1 or 1,000,000 times.

### 2.3 Source + block profiling

When designing our source level profiling system, we aimed to take advantage of prior work on low level profile directed optimizations. However, optimizations based on source-level profile information may result in a different set of blocks than the blocks generated for the profiled run of a program. If blocks are profiled naively, for instance, by assigning each block a number in the order in which the blocks are generated, then the block numbers will not be consistent after optimizing with source information. Therefore optimization using source profile information and those using block profile information cannot be done after a single profiled run of a program.

We take the naive approach to block profiling and use the following workflow to take advantage of both source and block level profile directed optimizations. First we compile and instrument a program to collect source-level information. We run this program and collect only source-level information. Next we recompile and optimize the program using the source-level information only, and instrument the program to collect block-level information. The profile directed meta-programs reoptimize at this point. We run this program and collect only the block-level information. Finally, we recompile the program with both source-level and block-level information. Since the source information has not changed, the meta-programs generate the same source code, and thus the compiler generates the same blocks. The blocks are then optimized with the correct profile information.

While the workflow seems to significantly complicate the compilation process, the different between using only block-level profiling and using both source-level and block-level profiling is small. To use any kind of profile directed optimizations requires a 300% increase in the number of steps (from compile to compile-profile-compile). To use both source-level and block-level profile directed optimizations requires only an additional 66% increase in number of steps (compile-profile-compile to compile-profile-compile-profile-compile).

## 3. Examples

This section presents several macros that use profiling information to optimize the expanded code. The first example demonstrates unrolling loops based on profile information. While loop unrolling can be done with low level profile information, we discuss when it can be useful or even necessary to do at the meta-programming level. The second example demonstrates call site optimization for a object-oriented DSL by reordering the clauses of a conditional branching structure, called `exclusive-cond`, based on profile information. The final example demonstrates specializing a data structure based on profile information.

### 3.1 Loop Unrolling

Loop unrolling is a standard compiler optimization. However, striking a balance between code growth and execution speed when unrolling loops is tricky. Profile information can help the compiler focus on the most executed loops.

Profile directed loop unrolling can be done using low-level profile information. However, loop unrolling at a low-level requires

associating loops with the low level profile structures, such internal nodes or even basic blocks, and cannot easily handle arbitrary recursive functions. More importantly, with the rise in interest and use of DSLs compilers, implementing loop unrolling via meta-programming may be necessary to get high performance loops in a DSL.

This loop example unrolls Scheme's named let<sup>1</sup>, as seen in figure 1. This defines a loop that runs for `i=5` to `i=0` computing factorial of 5. This named let might normally be implemented via a recursive function, as seen in figure 3. The example in figure 1 would produce a recursive function `fact`, and immediately call it on 5. With a reasonable compiler, this named let is equivalent to the C implementation in figure 2

```
(let fact ([i 5])
  (if (zero? i)
      1
      (* n (fact (sub1 n)))))
```

Figure 1: The most executed program in all of computer science

```
int i = 5;
int n = 1;
fact: if(i == 0){
  n;
} else {
  n = n * --i;
  goto fact;
}
```

Figure 2: And in C

```
(define-syntax let
  (syntax-rules ()
    [(_ name ([x e] ...) body1 body2 ...)
     ((letrec ([name (lambda (x ...) body1 body2 ...))
```

Figure 3: a simple definition of a named let

Figure 4 defines a macro, `named-let`, that unrolls the loop between 1 and 3 times, depending on profile information. At compile time, the macro-expander runs `(or (profile-query-weight #'b1) 0)`. This asks the runtime for the profile information associated with `b1`, the first expression in the body of the loop. Recall that `profile-query-weight` returns a value between 0 and 1 if there is profile information for a piece of syntax, and false otherwise. Using the profile weight, we calculate `unroll-limit`. If the profile weight is 1, meaning the expression is executed more than any other expression during the profiled run, `unroll-limit` is 3. If the weight is 0, meaning the expression is never executed during the profiled run, `unroll-limit` is 0. Finally, `named-let` generates a macro called `name`, where name is the identifier labeling the loop in the source code, does the work of unrolling the loop up to `unroll-limit` times.

In fact, a named let defines a recursive function and immediately calls it. While this can be used for simple loops, a named let

<sup>1</sup> Strictly speaking, we do not implement named let, since in loop unrolling macro, the name is not assignable.

```

(define-syntax named-let
  (lambda (x)
    (syntax-case x ()
      [(_ name ([x e] ...) b1 b2 ...)
       #'((letrec ([tmp (lambda (x ...)
                          #,(let* ([profile-weight
                                (or (profile-weight
                                     query-weight #'b1) 0)]
                                [unroll-limit
                                 (floor (* 3 profile-weight))])
                          #'(define-syntax name
                              (let ([count #, unroll-limit]
                                [weight #, profile-weight])
                                (lambda (q)
                                  (syntax-case q ()
                                    [(_ enew (... ..))
                                     (if (or (= count 0)
                                              (< weight 0)
                                              #'(tmp enew (... ..))
                                              (begin
                                                (set! count (+ count 1))
                                                #'((lambda (x ...)
                                                         enew (... ..))
                                                         b1 b2 ...))])
                                      tmp)
                                    e ...)))]))
        #'(define-syntax name
            (let ([count #, unroll-limit]
              [weight #, profile-weight])
              (lambda (q)
                (syntax-case q ()
                  [(_ enew (... ..))
                   (if (or (= count 0)
                           (< weight 0)
                           #'(tmp enew (... ..))
                           (begin
                             (set! count (+ count 1))
                             #'((lambda (x ...)
                                    enew (... ..))
                                    b1 b2 ...))])
                     tmp)
                  e ...)))]))

```

Figure 4: a macro that does profile directed loop unrolling

may have non-tail calls or even multiple recursive calls along different branches. This macro does more than loop unrolling—it does recursive function lining. A more clever macro could unroll each call site a different number of times, depending on how many times that particular call is executed. This would allow more fine grain control over code growth. For brevity, we restrict the example and assume `named-let` is used as a simple loop. Each call site is unrolled the same number of times.

Similar macros are easy to write for `do` loops, and even for `letrec` to inline general recursive functions.

### 3.2 exclusive-cond

In this section we present a branching construct called `exclusive-cond` that can automatically reorder the clauses based on which is mostly likely to be executed. This optimization is analogous to basic block reordering, but operates at a much higher level.

We consider this construct in the context of an object-oriented DSL with classes, inheritance, and virtual methods, similar to C++. Consider a class with a virtual method `get_x`, called `Point`. `CartesianPoint` and `PolarPoint` inherit `Point` and implement the virtual `get_x`. We will use `exclusive-cond` to inline virtual method calls.

`cond` is a Scheme branching construct analogous to a series of `if` or `if-else` statements. The clauses of `cond` are executed in order until the left-hand side of a clause is true. If there is an `else` clause, the right-hand side of the `else` clause is taken only if no other clause's left-hand side is true.

Figure 5 shows an example of a `cond` generated by our hypothetical OO DSL. The DSL compiler simply expands every virtual

method call into a conditional branch for known instances of an object.

```

(cond
  [(class-equal? obj CartesianPoint)
   (field obj x)]
  [(class-equal? obj PolarPoint)
   (* (field obj rho) (cos (field obj theta)))]
  [else (method obj "get_x")])

```

Figure 5: An example of `cond`

By profiling the branches of the `cond`, we can sort the clauses in order of most likely to succeed, or even drop clauses that occur too infrequently inline. However, `cond` is order dependent. While the programmer can see the clauses are mutually exclusive, the compiler cannot prove this in general and cannot reorder the clauses.

Instead of wishing our compiler was more clever, we use meta-programming to take advantage of this high-level knowledge. We define `exclusive-cond`, figure 6, with the same syntax and semantics of `cond`<sup>2</sup>, but with the restriction that clause order is not guaranteed. We then use profile information to reorder the clauses.

The `exclusive-cond` macro will rearrange clauses based on the profiling information of the right-hand sides. Since the left-hand sides will be executed depending on the order of the clauses, profiling information from the left-hand side is not enough to determine which clause is true most often.<sup>3</sup> The clause record stores the original syntax for the clause and the weighted profile count for that clause. Since a valid `exclusive-cond` clause is also a valid `cond` clause, the syntax is simply copied, and a new `cond` is generated with the clauses sorted according to profile weights. If an `else` clause exists then it is emitted as the final clause.

Figure 7 shows an example of `exclusive-cond` and the code to which it expands. In this example, we assume the object is a `PolarPoint` most of the time.

#### 3.2.1 Another use of `exclusive-cond`

`case` is a pattern matching construct that is easily given profile directed optimization by implementing it in terms of `exclusive-cond`. `case` takes an expression `key-expr` and an arbitrary number of clauses, followed by an optional `else` clause. The left-hand side of each clause is a list of constants. `case` executes the right-hand side of the first clause in which `key-expr` is `eqv?` to some element of the left-hand. If `key-expr` is not `eqv?` to any element of any left-hand side and an `else` clause exists then the right-hand side of the `else` clause is executed.

Figure 8 shows an example `case` expression. If `x` is 1, 2, or 3, then `e1` is executed. If `x` is 4 or 5, then `e2` is executed. Note that while 3 appears in the second clause, if `x` is 3 then `e1` will be evaluated. The first occurrence always take precedence.

Since `case` permits clauses to have overlapping elements and uses order to determine which branch to take, we must remove overlapping elements before clauses can be reordered. Each clause is parsed into the set of left-hand side keys and right-hand side bodies. Overlapping keys are removed by keeping only the first instance of each key when processing the clauses in the original order. After removing overlapping keys, an `exclusive-cond` is generated.

<sup>2</sup>We omit the alternative `cond` syntaxes for brevity.

<sup>3</sup>Schemers will note this means we cannot handle the single expression `cond` clause syntax.

```

(define-syntax exclusive-cond
  (lambda (x)
    (define-record-type clause (fields syn weight))
    (define (parse-clause clause)
      (syntax-case clause ()
        [(e0 e1 e2 ...) (make-
          clause clause (or (profile-query-
            weight #'e1) 0))])
        [_ (syntax-error clause "invalid
          clause")]))
    (define (sort-clauses clause*)
      (sort (lambda (cl1 cl2)
              (> (clause-
                weight cl1) (clause-weight cl2)))
            (map parse-clause clause*)))
    (define (reorder-cond clause* els)
      #'(cond
        #,@(map clause-syn (sort-
          clauses clause*))
        #,@(if els #'(,els) #'())))
    (syntax-case x (else)
      [( _ m1 ... (else e1 e2 ...) ) (reorder-
        cond #'(m1 ...) #'(else e1 e2 ...))]
      [( _ m1 ... ) (reorder-
        cond #'(m1 ...) #f)])))

```

Figure 6: Implementation of `exclusive-cond`

```

(exclusive-cond
  [(class-equal? obj CartesianPoint) (field obj x)
   executed 2 times]
  [(class-equal? obj PolarPoint)
   (* (field obj rho) (cos (field obj theta)))]
   executed 5 times]
  [else (method obj "get_x")] ; executed
  8 times

(cond
  [(class-equal? obj PolarPoint) (* (field obj rho)
   (cos (field obj theta)))]
  [(class-equal? obj CartesianPoint) (field obj x)]
  [else (method obj "get_x")] ; executed
  8 times.

```

Figure 7: An example of `exclusive-cond` and its expansion

```

(case x
  [(1 2 3) e1]
  [(3 4 5) e2]
  [else e3])

```

Figure 8: An example of a `case` expression

Figure 9 shows how the example `case` expression from figure 8 expands into `exclusive-cond`. Note the duplicate 3 in the second clause is dropped to preserve ordering constraints from `case`.

```

(exclusive-cond x
  [(memv x (1 2 3)) e1]
  [(memv x (4 5)) e2]
  [else e3])

```

Figure 9: The expansion of figure 8

```

(define-sequence-datatype seq1 (0 3 2 5)
  seq? seq-map seq-first seq-ref seq-
  set!)

```

Figure 11: Use of the `define-sequence-datatype` macro

### 3.3 Data type Selection

The previous examples show that we can easily bring well-known optimizations up to the meta-level, enabling the DSL writer to take advantage of traditional profile directed optimizations. While profile directed meta-programming enables such traditional optimizations, it also enables higher level decisions normally done by the programmer.

In this example we present a library that provides a sequence datatype. We consider this in the context of a DSL or library writer whose users are domain experts, but not computer scientists. While a domain expert writing a program may know they need a sequence for their program, they may not have the knowledge to figure out if they should use a tree, or a list, or a vector. Past work has bridged this gap in knowledge by providing tools that can recommend changes and provide feedback. We take this a step further and provide a library that will automatically specialize the data structure based on usage.

The example in figure 10 chooses between a list and a vector using profile information. If the program uses `seq-set!` and `seq-ref` operations more often than `seq-map` and `seq-first`, then the sequence is implemented using a `vector`, otherwise using a `list`.

Figure 11 demonstrates the usage of the `define-sequence-datatype` macro. In this example, a sequence named `seq1` is defined and initialized to contain elements 0, 3, 2, and 5. The macro also takes the various sequence operations as arguments, though this is a hack. To get unique per sequence source information, we simply use the source information from those extra arguments. A production example would omit this hack.

The macro expands into a series of definitions for each sequence operation and a definition for the sequence datatype. This example redefines the operations for each new sequence and evaluates the name to ensure function inlining does not distort profile counts. A clever compiler might try to throw out the effect-free reference to `name` in the body of each operation, so this implementation is fragile.

## 4. Implementation

This section describes our implementation of the profiling system, and how source-level and block-level profile directed optimizations can work together in our system. First we present how code is instrumented to collect profile information. Then we present how profile information is stored and accessed. Finally we present how we use both source-level and block-level profile directed optimizations in the same system.

To hell with this example. We need to break it up and make it slightly more sensible to the source information, but that would make this hacking easier. Maybe move choose

nonsense to an appendix and just focus on the macro itself. I'm going to need Kent to check this section.

```

(define-syntax define-sequence-datatype
  (let ([ht (make-eq-hashtable)])
    (define args
      `((seq? . #'(x))
        (seq-map . #'(f s))
        (seq-first . #'(s))
        (seq-ref . #'(s n))
        (seq-set! . #'(s i obj))))
    (define defs
      `((make-seq      ,#'list . ,#'vector)
        (seq?         ,#'list? . ,#'vector?)
        (seq-map      ,#'map . ,#'for-each)
        (seq-first    ,#'car . ,#' (lambda (x) (vector-ref x 0)))
        (seq-ref      ,#'list-ref . ,#'vector-ref)
        (seq-set!     ,#' (lambda (ls n obj) (set-car! (list-tail ls n) obj)) . ,#'vector-
set!)))
    (define (choose-args name)
      (cond
        [(assq name args) => cdr]
        [else (syntax-error name "invalid method:")]))
    (define (choose name)
      (let ([seq-set!-count (hashtable-ref ht 'seq-set! 0)]
            [seq-ref-count (hashtable-ref ht 'seq-ref 0)]
            [seq-first-count (hashtable-ref ht 'seq-first 0)]
            [seq-map-count (hashtable-ref ht 'seq-map 0)])
        (cond
          [(assq name defs) =>
            (lambda (x)
              (let ([x (cdr x)])
                (if (> (+ seq-set!-count seq-ref-count)
                      (+ seq-first-count seq-map-count))
                    (cdr x)
                    (car x)))))]
          [else (syntax-error name "invalid method:")])])
    (lambda (x)
      (syntax-case x ()
        [(_ var (init* ...) name* ...)
         (for-each
          (lambda (name)
            (hashtable-set! ht name
              (or (profile-query-weight name) 0)))
            (map syntax->datum #'(name* ...)))
          (with-syntax [(body* ...) (map (lambda (name) (choose (syntax->datum name))) #'(name* ...)))]
            [(args* ...) (map (lambda (args) (choose-args (syntax->datum name))) #'(name* ...)))]
            #'(begin (define (name* args* ...) (begin name* (body* args* ...))) ...
              (define var (#, (choose 'make-seq init* ...)))))))]))

```

---

Figure 10: a macro that defines a sequence datatype based on profile information



## 4.1 Instrumenting code

The naive method for instrumenting code to collect source profile information is to attach the source information to each AST node internally. At an appropriately low level, that source information can be used to generate code that increments profile counters. However this method can easily distort the profile counts. As nodes are duplicated or thrown out during optimizations, the source information is also duplicated or lost.

Instead we create a separate profile form that is created during macro expansion. Each expression `e` that has source information attached is expanded internally to `(begin (profile src) e)`, where `src` is the source object attached to `e`. The profile form is considered an effectful expression internally and should never be thrown out or duplicated, even if `e` is.

These profile forms are retained until basic blocks are generated. While generating basic blocks, the source objects from the profile forms are gathered up and attached to the basic block in which they appear. When a basic-block is entered, every instruction in that block will be executed, so any profile counters in the block must be incremented. Since all the profile counters must be incremented, it is safe to increment them all at the top of the block.

In our implementation, we attempt to minimize the number of counters executed at runtime. After generating basic blocks and attaching the source objects to their blocks, we analyze the blocks to determine which counters can be calculated in terms of other counters. If possible, a counter is computed as the sum of a list of counters (+counters) minus the sum of a list of counters (-counters). This complicated the internal representation of counters and the generation of counters, but decreases the overhead of profiling.

To instrument block-level profiling, we reuse the above infrastructure by creating fake source objects. When a file is compiled, we reset global initial block number to 0, and create a fake source file descriptor based on the file name. When creating blocks, each block is given a source object using the fake file descriptor, and using the blocks number as the starting and ending file position. This fake source object is used when block-level profiling is enable. This fake source is ignored and the list of sources from the source code is used when source-level profiling is enable.

## 4.2 Storing and Loading profile data

We store profile data by creating a hash table from source file names to hash tables. Each second level hash table maps the starting file position of the expression to the weighted count of the expression. This lookup table is only populated after loading profile data from a file and not from a current profiled run. After loading profile data, it is accessible through `profile-query-weight`.

Profile data is not immediately loaded into the lookup table after a profiled run of a program. Profile data must first be dumped via `profile-dump-data` and then loaded via `profile-load-data`.

To dump profile data, the run time gathers up all profile counters. Recall that some counters are computed indirectly in terms of other counters. The values for these indirect counters are computed. These values with their associated source objects are then written to a file.

To support loading multiple data sets, we do not load execution counts directly into the lookup table. Instead we compute the percent of max for each counter. Before loading a new data set, we find the maximum counter value. Each weighted count is computed as a percent of the maximum counter value. If an entry for a source already exists in the lookup table then we compute the weighted average of the previous entry and the counter we're currently loading. We store the weighted count and the current weight in the lookup table, incrementing the weight by one with each new data set.

## 5. Related and Future Work

Modern systems such as GCC, .NET, and LLVM use profile directed optimizations [3, 4, 5]. However, these systems provide mostly low level optimizations, such as optimizations for block order and register allocation. In addition to limiting the kinds of optimizations the compiler can do, this low-level profile information is fragile.

Recently there has been work to give programmers advice on which data structure to use <http://dx.doi.org/10.1109/CGO.2009.36>, but with our techniques we can automatically optimize the generated code instead of just advice the programmer.

GCC profiles an internal control-flow graph (CFG). To maintain a consistent CFGs across instrumented and optimization builds, GCC requires similar optimization decisions across builds. By associating profile information with source expression we can more easily reuse profile information [1]. In our system, all profile information for a source file is usable as long as the source file does not change.

.NET provides some higher level optimizations, such as function inlining and conditional branch optimization similar to `exclusive-cond` and `case` presented here. To optimize `switch` statements, .NET uses *value* profiling in addition to execution count profiling [5]. By probing the values used in a switch statement, the compiler can attempt to reorder the cases of the `switch` statement.

The standard model for profile directed optimizations requires probes the instrument-profile-optimize workflow. LLVM has a different model for profile directed optimization. LLVM uses a runtime re-optimizer that monitors the running program. The runtime reoptimizer can profile the program as it runs "in the field" and perform simple optimizations to the machine code, or call off to an offline optimizer for more complex optimizations on the LLVM bytecode.

Meta-programs generate code at compile time, so the examples presented in section 3 require the standard instrument-profile-optimize workflow. However, because we expose an API to access profiling information, we could use this system to perform runtime decisions based on profile information. To truly be beneficial, this requires keeping the runtime overhead of profiling very low, which is not usually the case [1, 2]. However, our techniques for reducing the number of counters and our careful representation of profile forms allows accurate source profiling with little overhead.

## Bibliography

- [1] Deheo Chen, Neil Vachharajani, Robert Hundt, Shih-wei Liao, Vinodha Ramasamy, Paul Yuan, Wenguang Chen, and Weimin Zheng. Taming Hardware Event Samples for FDO Compilation. In *Proc. Annual IEEE/ACM international symposium on Code generation and optimization*, 8, pp. 42–52, 2010. [http://hpc.cs.tsinghua.edu.cn/research/cluster/papers\\_cwg/tamingsample.pdf](http://hpc.cs.tsinghua.edu.cn/research/cluster/papers_cwg/tamingsample.pdf)
- [2] Thomas M Conte, Kishore N Menezes, and Mary Ann Hirsch. Accurate and practical profile-driven compilation using the profile buffer. In *Proc. Annual ACM/IEEE international symposium on Microarchitecture*, 29, pp. 36–45, 1996. [http://pdf.aminer.org/000/244/348/commercializing\\_profile\\_driven\\_optimization.pdf](http://pdf.aminer.org/000/244/348/commercializing_profile_driven_optimization.pdf)
- [3] LLVM: An infrastructure for multi-stage optimization. Master dissertation, University of Illinois, 2002.
- [4] Optimize Options - Using the GNU Compiler Collection. 2013.

I'm not sure what I'm doing with this section yet.

Value probes seem like a pretty ad-hoc method to get a very specific optimization. I don't know if I want to say that. measure overhead on a standard set of benchmarks. The benchmarks I ran at cisco suggest ~10% overhead, but those are not publicly accessible. This sentence belongs in implementation

- [5] Profile-Guided Optimizations. 2013. [http://msdn.microsoft.com/en-us/library/e7k32f4k\(v=vs.90\).aspx](http://msdn.microsoft.com/en-us/library/e7k32f4k(v=vs.90).aspx)